


Team Pineapple: Veritas

By Nicholas Ching, Suwei Ma, Avery Fan, Naren Loganathan

AI-Generated Detection Service Demo



Agenda

1

The problem we're trying to solve

2

Our solution

3

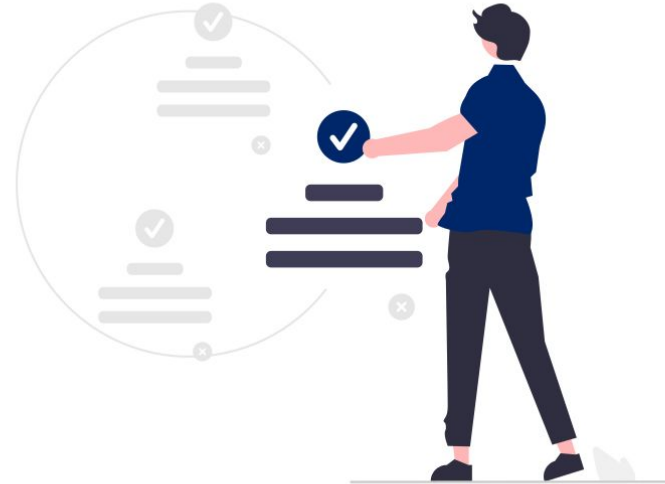
A demo

4

Why microservice architecture?

5

Future development goals



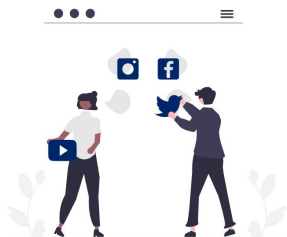


Discourse is the foundation of society.

An existential threat to society

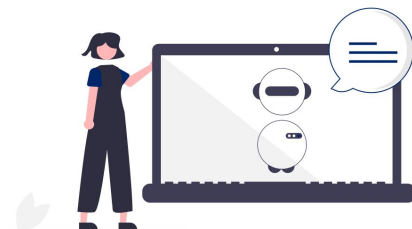
Social media and AI have fundamentally changed the way we engage with one another

How we engage



- Vastly greater platform to discuss and share ideas
- Greater threat that bad actors get **amplified**

The rise of AI



- AI has made bots harder to detect
- Easy to plug and play ChatGPT
- Decreased bot detectors **effectiveness by 30%**

The threat to society

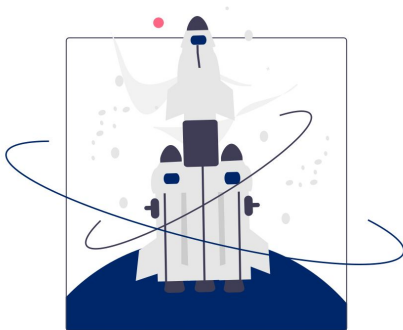


- In 2024 UK election, 45 bots published 450,000 tweets that got **3B impressions**
- During Trump's impeachment, bots accounted for 1% of users, but accounted for **31% of posts**

An Overview of Veritas

Our goal is to do one thing, but try and do it well. It's not sexy and not perfect. But that's the point.

What is our job?



Detect AI-Generated Content

By doing one thing well, our service can fit into any organizations workflow

What does our service enable?

Bot monitoring

Trust

Plagiarism checkers

Discourse

Plug & Play

Fact Checking

It is not our responsibility to tell the user what to do with the content

How our service integrates with clients

Create an incredibly simple interface

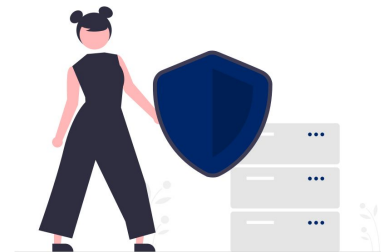


Client

Some text



A boolean + metadata



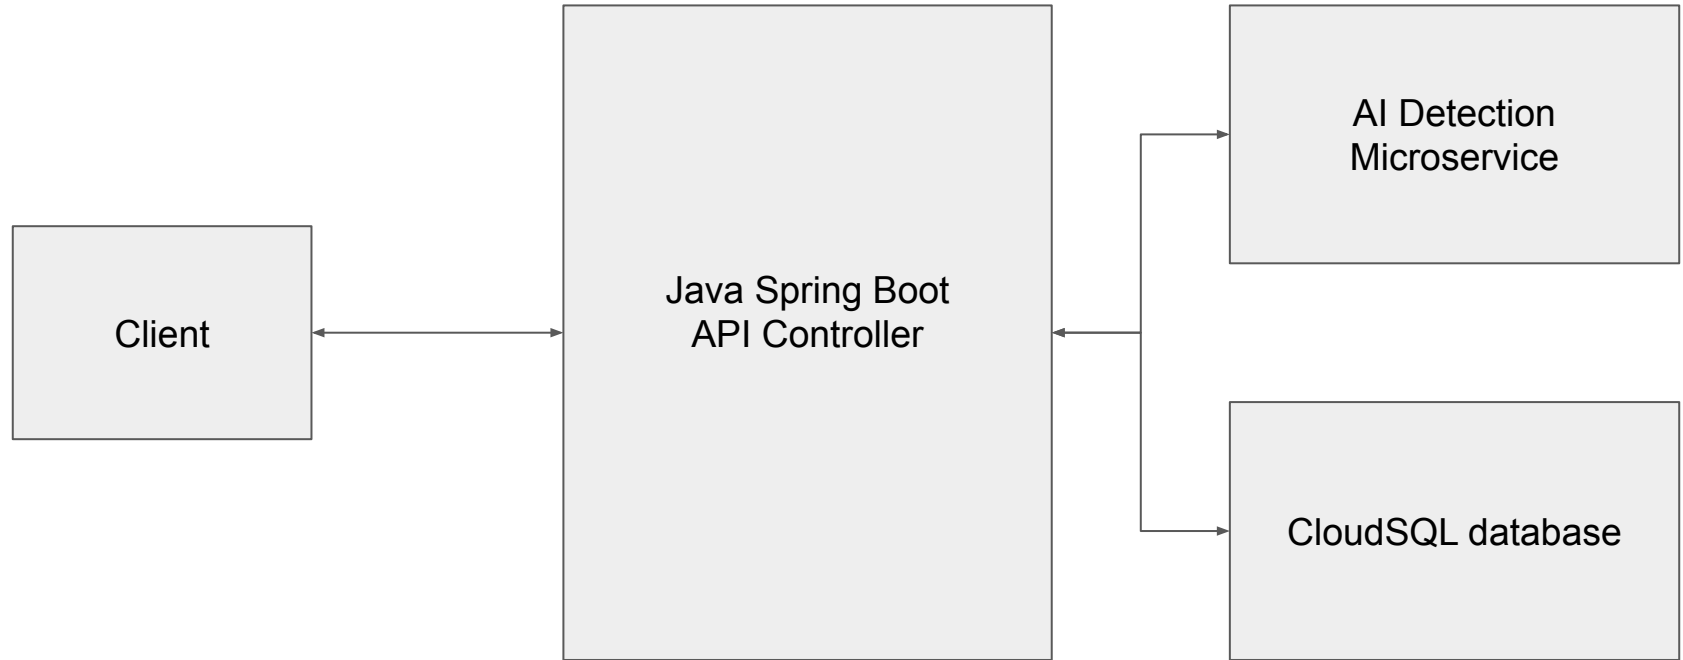
Veritas API

A demo of our service



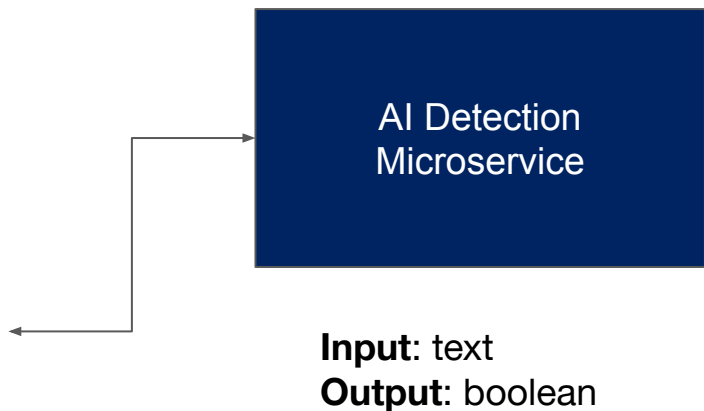
Our Architecture

A modular approach to AI-generated text detection



Why Microservice?

Modularity allows us to loosely couple our logic. We can iterate much faster this way



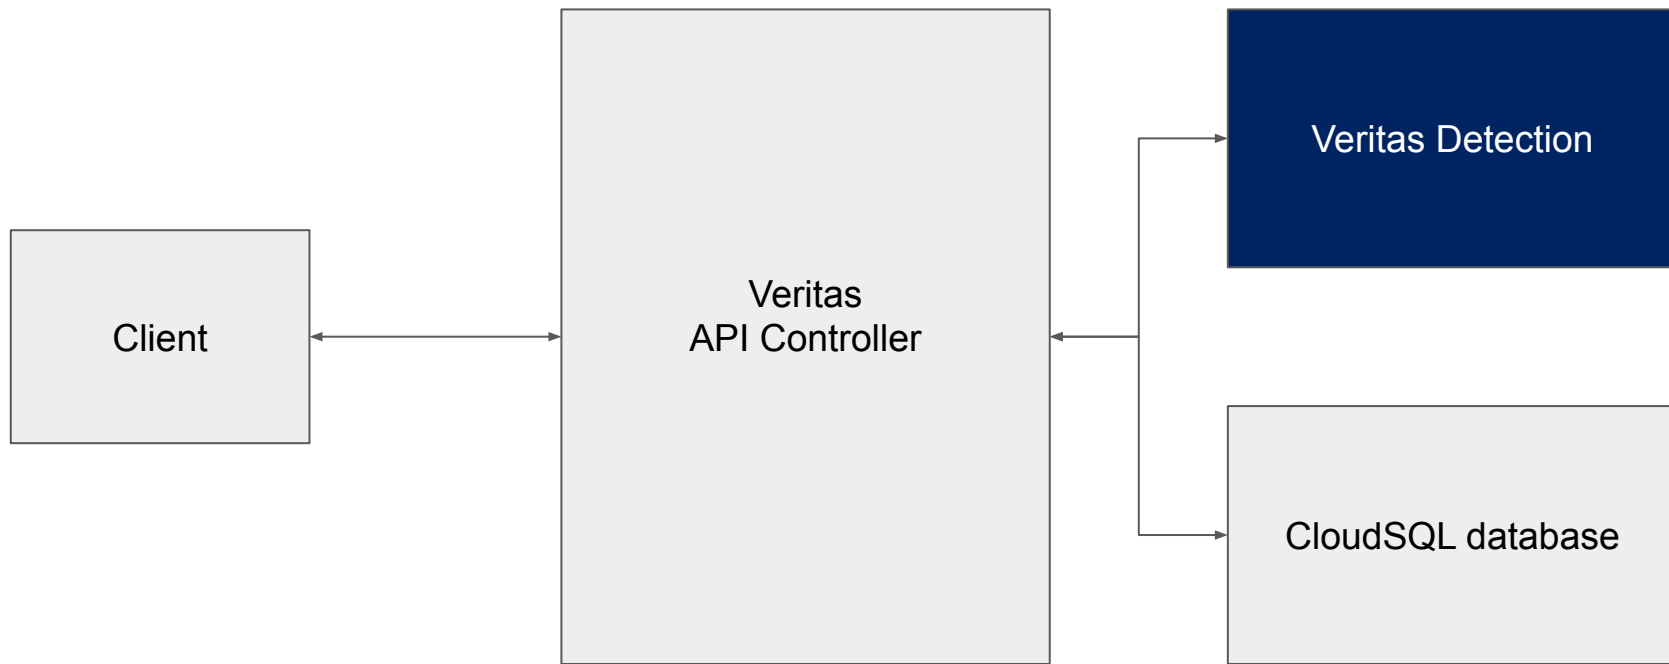
Focus on **What** not **How**

There are **no dependencies** between our main Java API controller and the detection system

We can debug, fine-tune, iterate and deploy all without needing to change our main API controller.

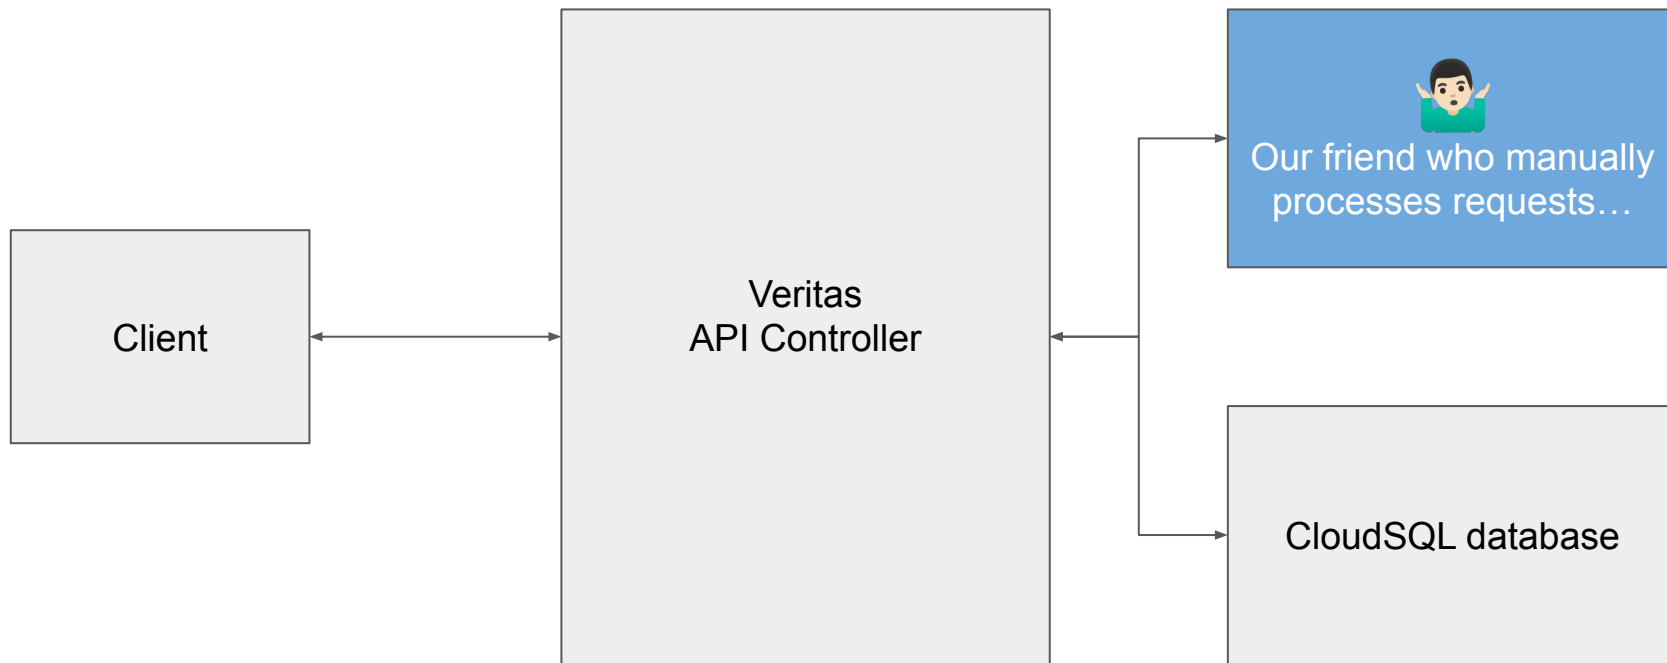
Our Architecture

We can plug and play here...



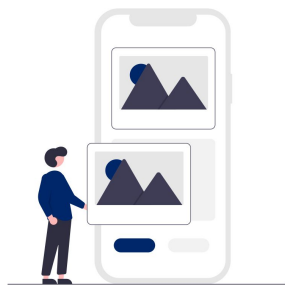
Our Architecture

We can plug and play here...

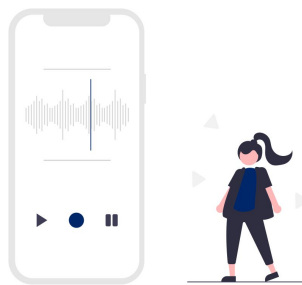


Future development goals

The goal is to provide a flexible service that can detect ALL generated content



Images



Audio



Text

There doesn't currently exist a foolproof way to detect if text, images or audio is generated by AI. So we iterate, and we keep up with the latest literature



Questions?